



Los Alamos NATIONAL LABORATORY

EST. 1943

Next-generation cluster management architecture

Towards Continuous Uptime



Paul Peltz Jr. and Lowell Wofford

11/11/2018



Operated by Los Alamos National Security, LLC for the U.S. Department of Energy's NNSA

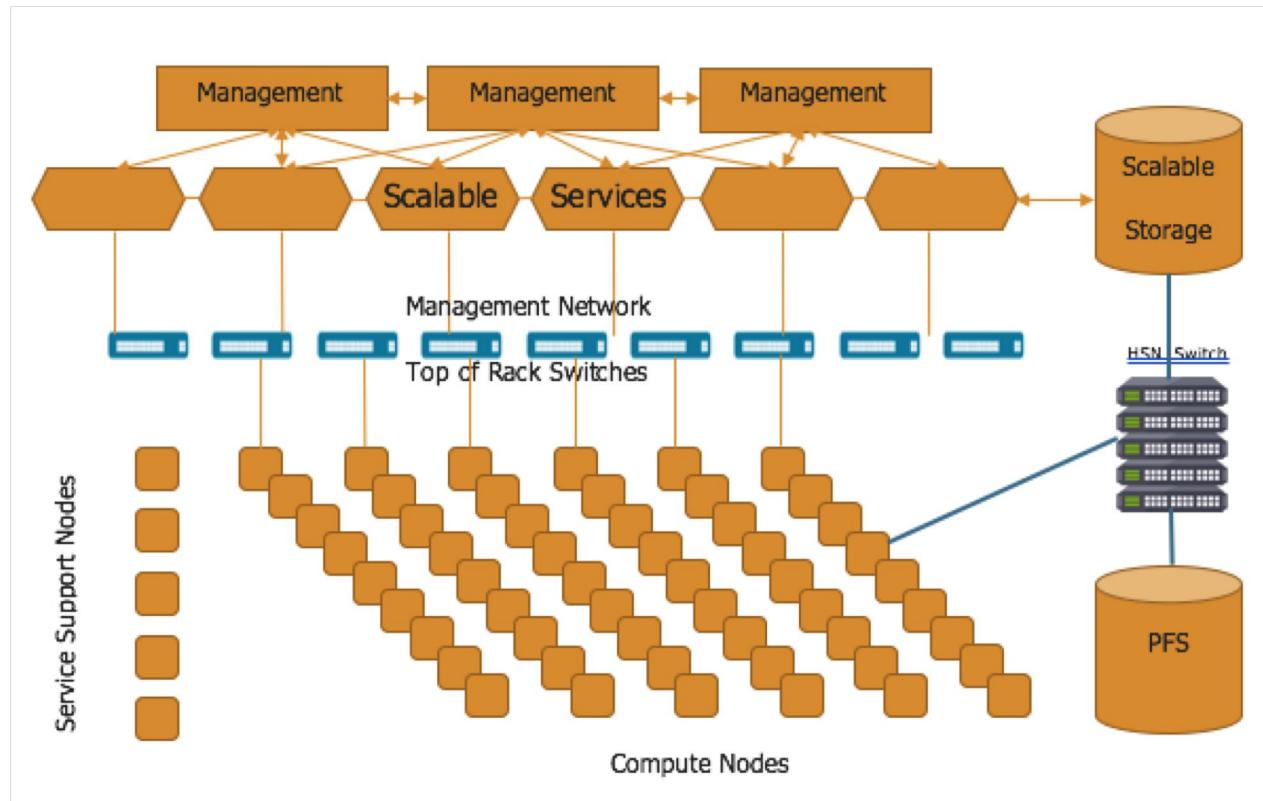
Current Issues

- **Current Cluster Management Tools**
 - Centralized
 - Monolithic
 - Proprietary/Outdated
 - Not Configuration Management and Orchestration capable
- **Downtime**
 - Binary Network Topologies (Up/Down)
 - Updates/Patches
- **Schedulers**
 - Lack integration into the cluster management
 - Inference model only

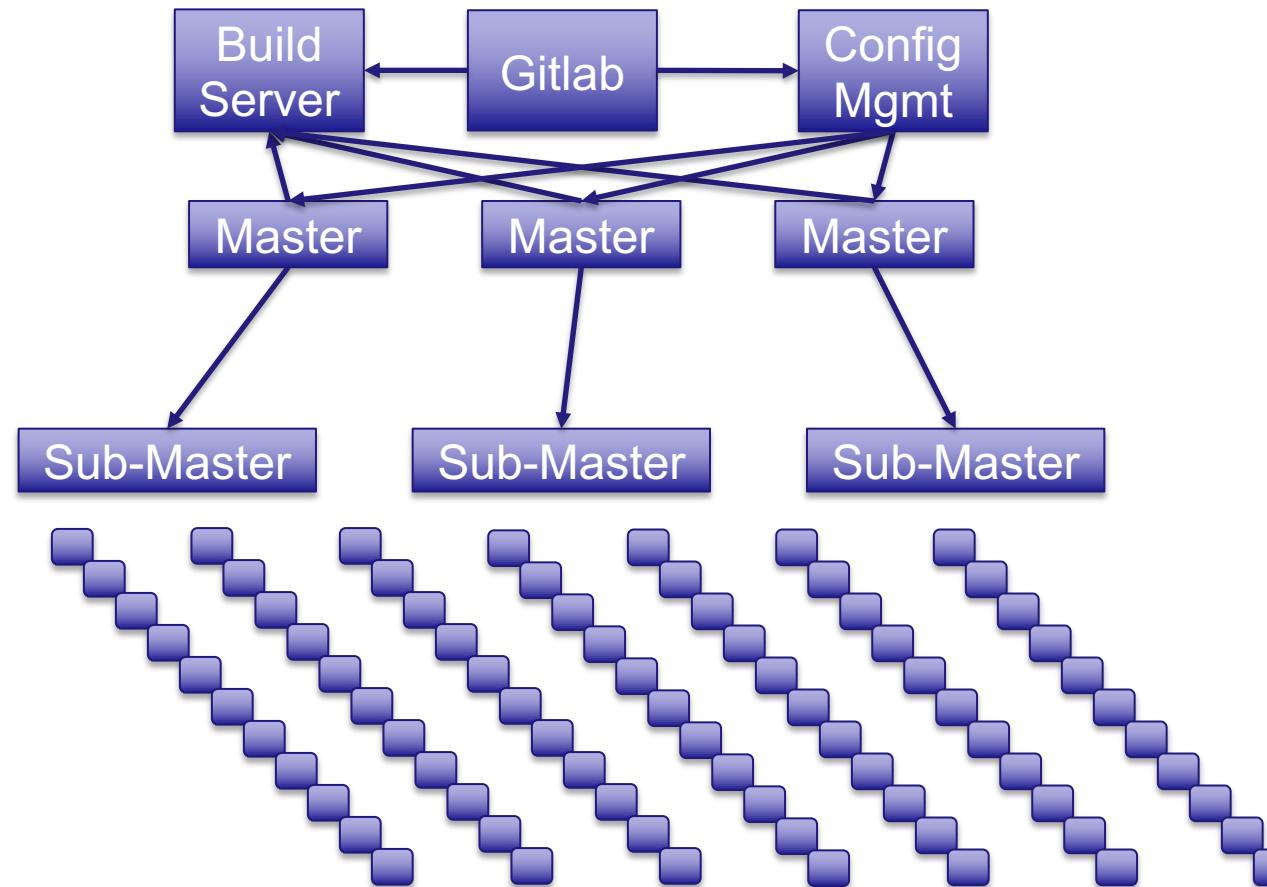
How do we solve these problems?

- **Current Cluster Management Tools**
 - Complete redesign necessary
 - Distributed State Engine, Modular
- **Downtime**
 - System should never go down
 - Degraded modes
 - Distributed Services
 - Rolling Upgrades
 - CI/CD of Patches and Changes
- **Schedulers**
 - Ties directly into the distributed state engine

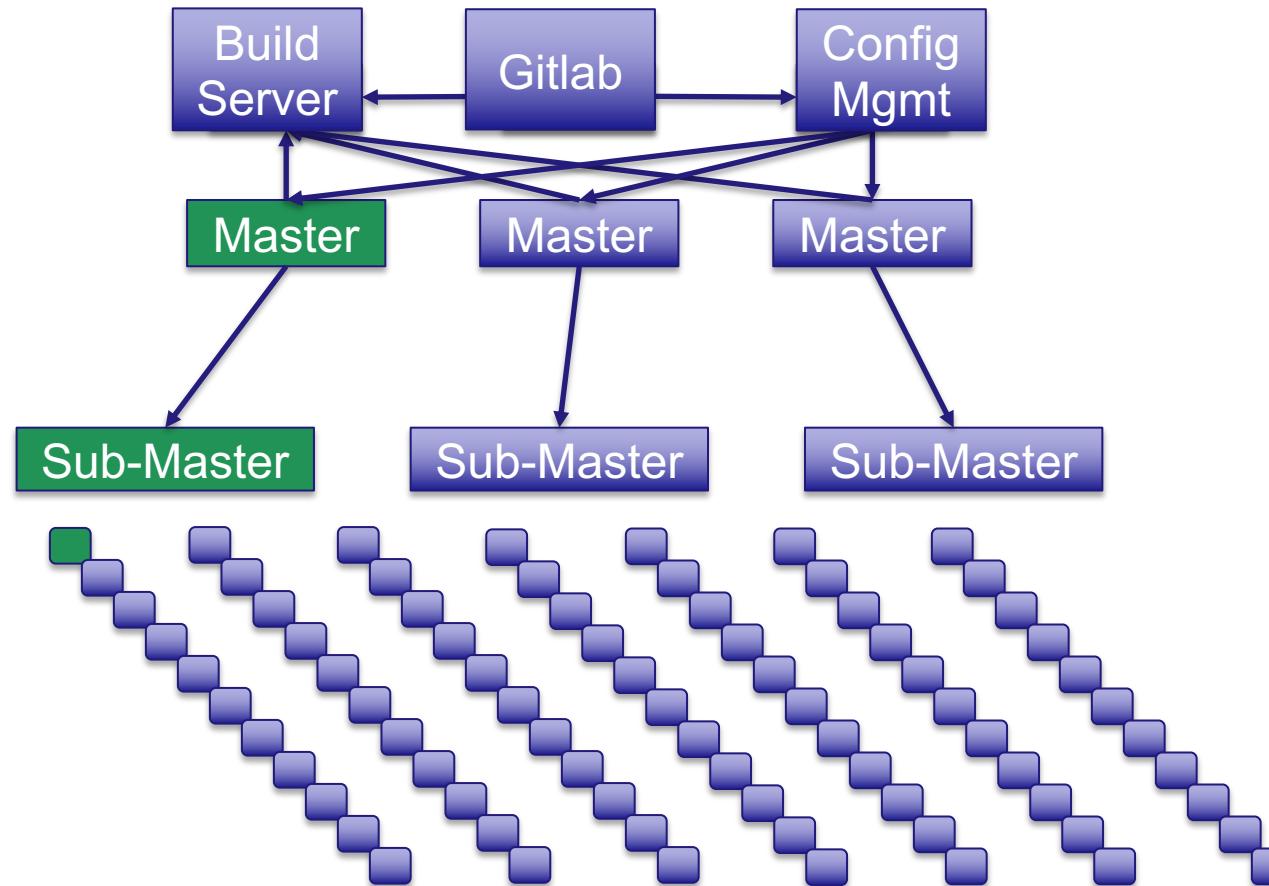
Redesign of HPC Architecture



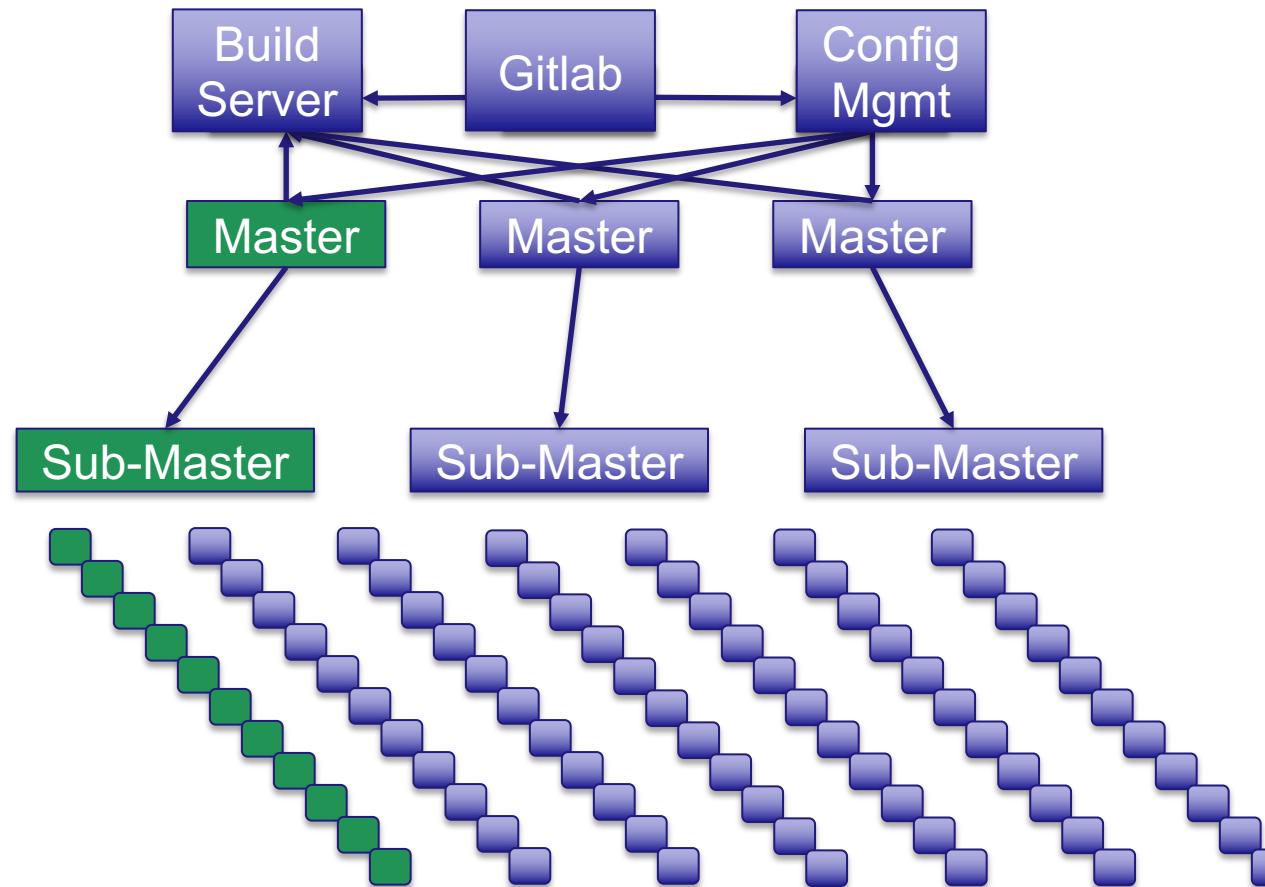
How we will test and deploy changes in the future



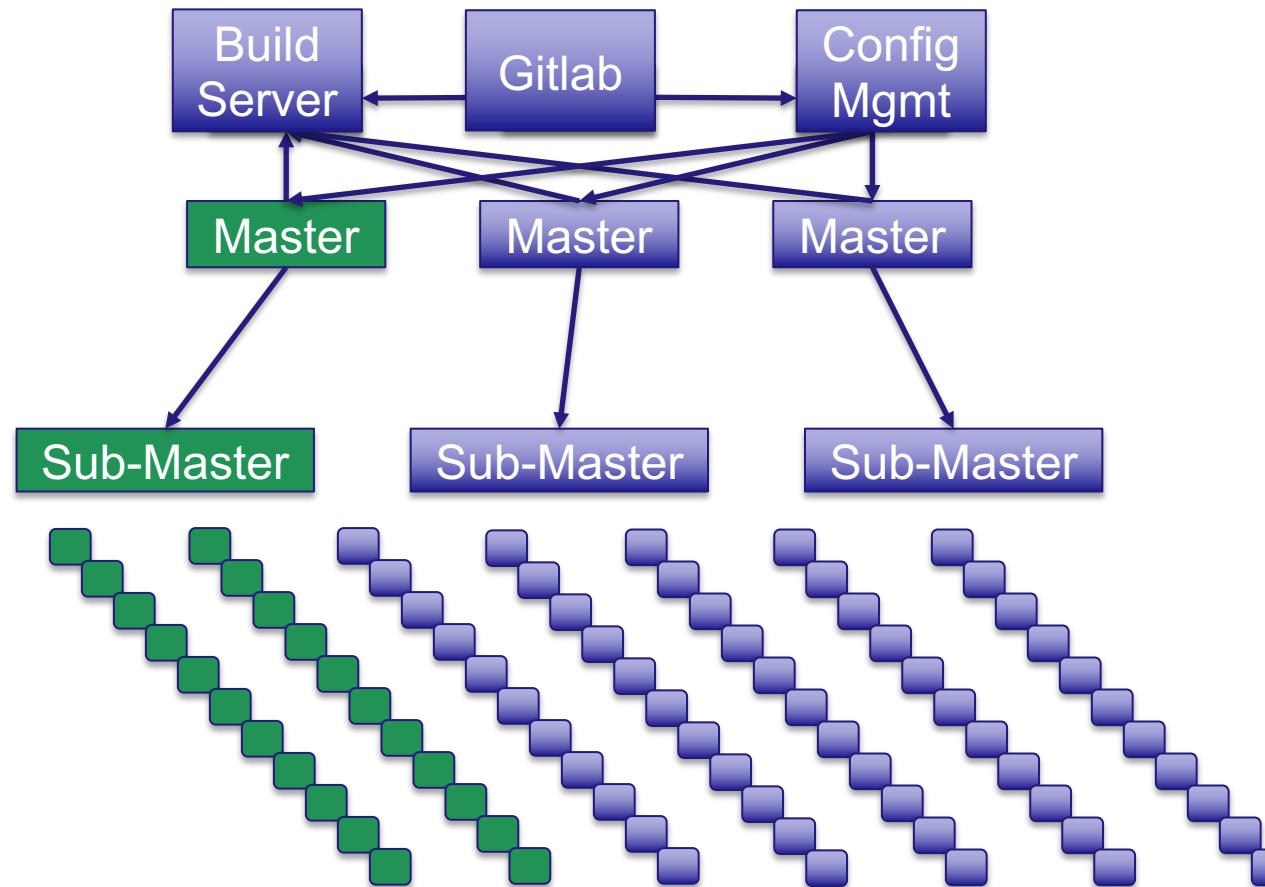
Stage 1 – Functional Testing



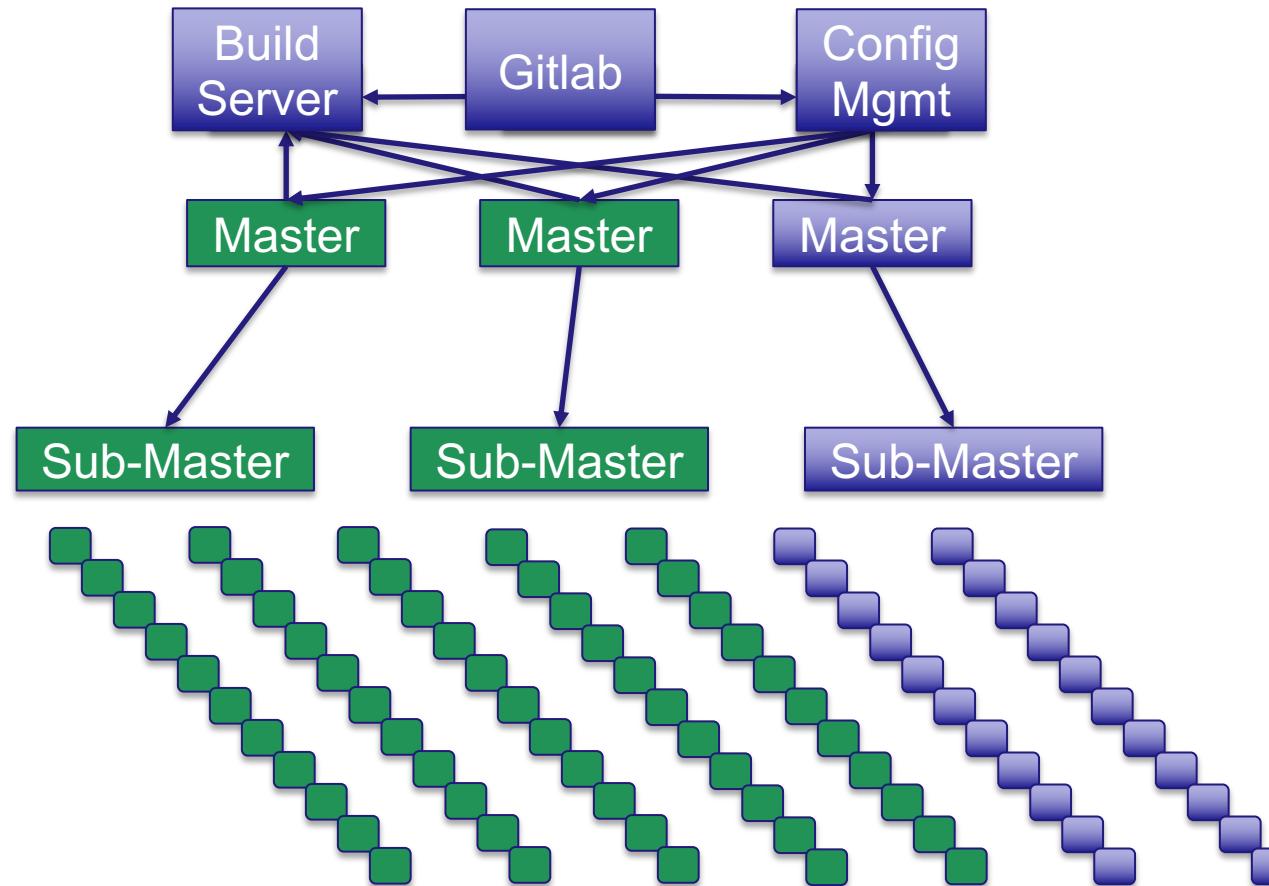
Stage 2 – Small Scale Test



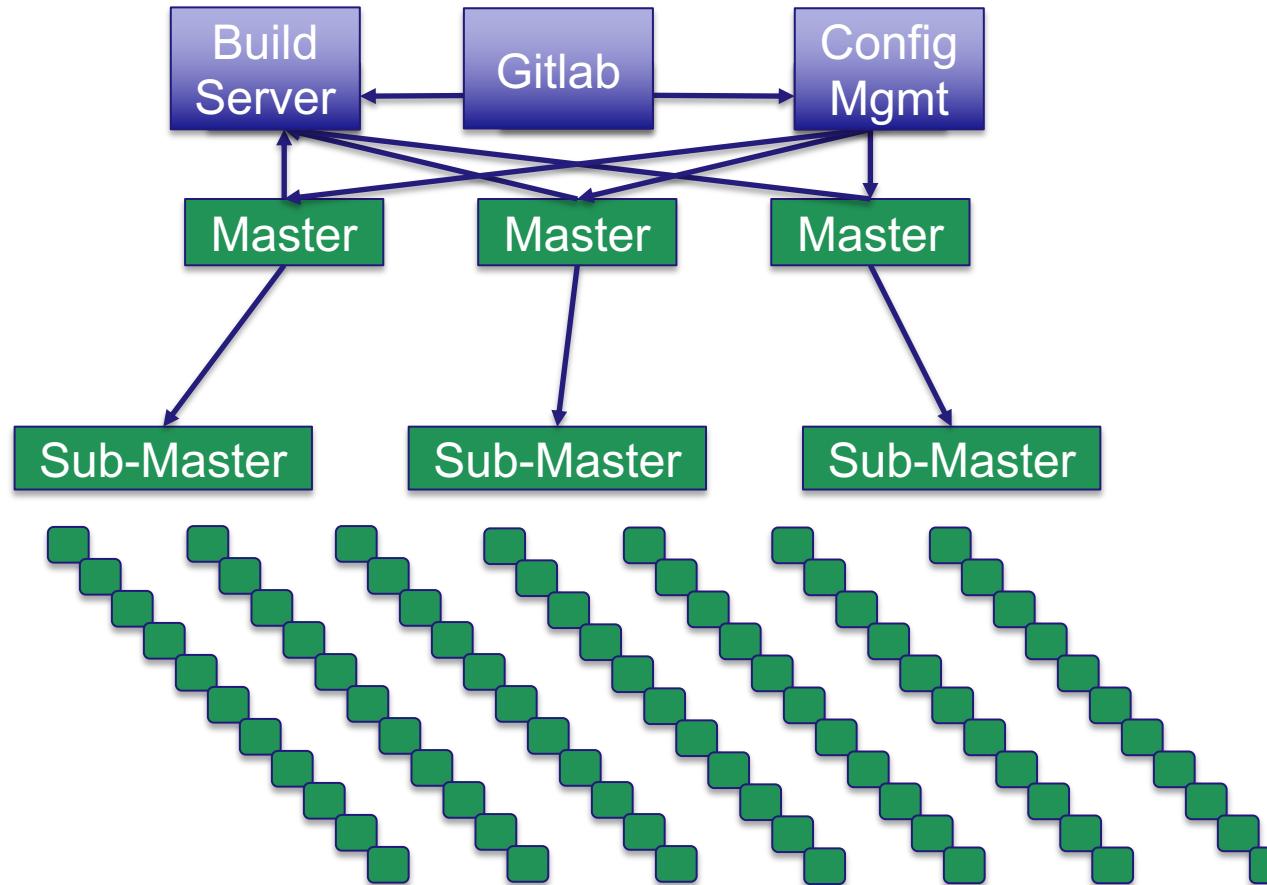
Stage 3 Scale Testing



Stage 4 – Update Additional Master/Sub-Master



Final Stage – Continue to reboot nodes into new images once jobs complete



Sounds great, how do we do this?

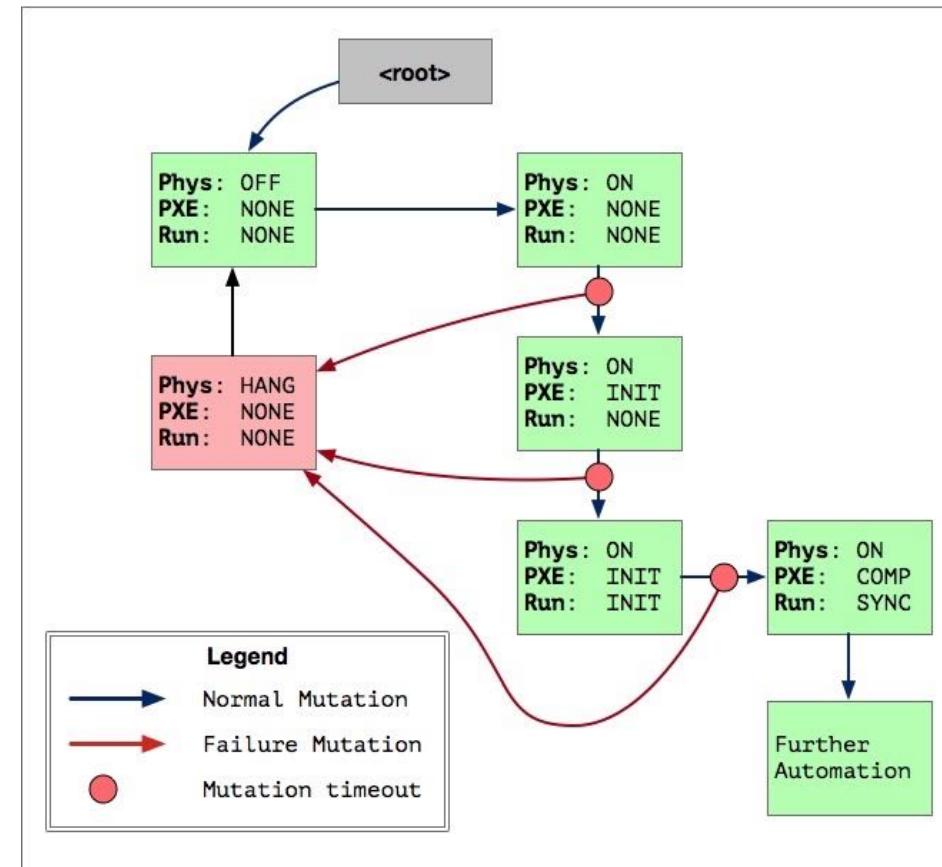
- **Service decentralization**
 - HA is not acceptable
 - VMs/Micro Services
 - Migrate services
- **Distributed State Engine**
 - The cluster manager must be able to track state, transitions, and final states of the system including the service nodes.
- **Backwards compatible APIs**
 - Need at least one version compatibility between service APIs.
- **Scheduler Integration**
 - Must be able to schedule the transition of release states on the system
 - Handle A/B configuration if there are compatibility issues

Automated and Continuous Testing!

- **Continuous Testing**
 - Testing first on development systems (TDS)
 - Automated testing for functionality and performance during each stage of deployment
 - Once release is approved that branch is released into production
 - Configuration management notifies state engine to begin rolling through upgrade
 - Testing is performed for each step
 - If a test exception occurs, human intervention will be required to resolve the issue
 - Roll forward or roll back support

Kraken is a State Engine for System Automation

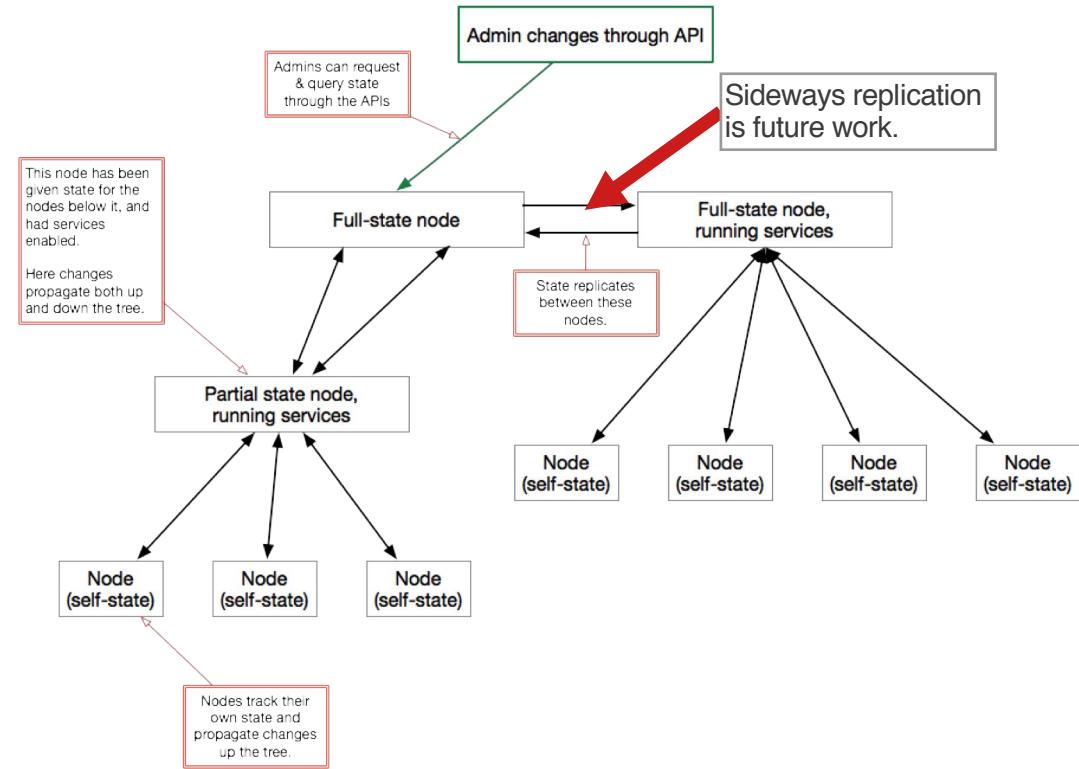
- **It can:** boot systems, load images, recover from failures... and more.
- Kraken tracks two kinds of **state**:
 - *What state are we in?* (“Discoverable”)
 - *What state should we be in?* (“Configuration”)
- Kraken knows how to **mutate** from the state you’re in to the state you want:
 - Maintains a directed graph of known state *mutations* (e.g. *power_off* -> *power_on*)
 - Searches the graph to create a chain of mutations to get from where we are to where we want to be.



Kraken is Distributed

Kraken uses a *protobuf*-based, light-weight datagram protocol to **synchronize** state across the system.

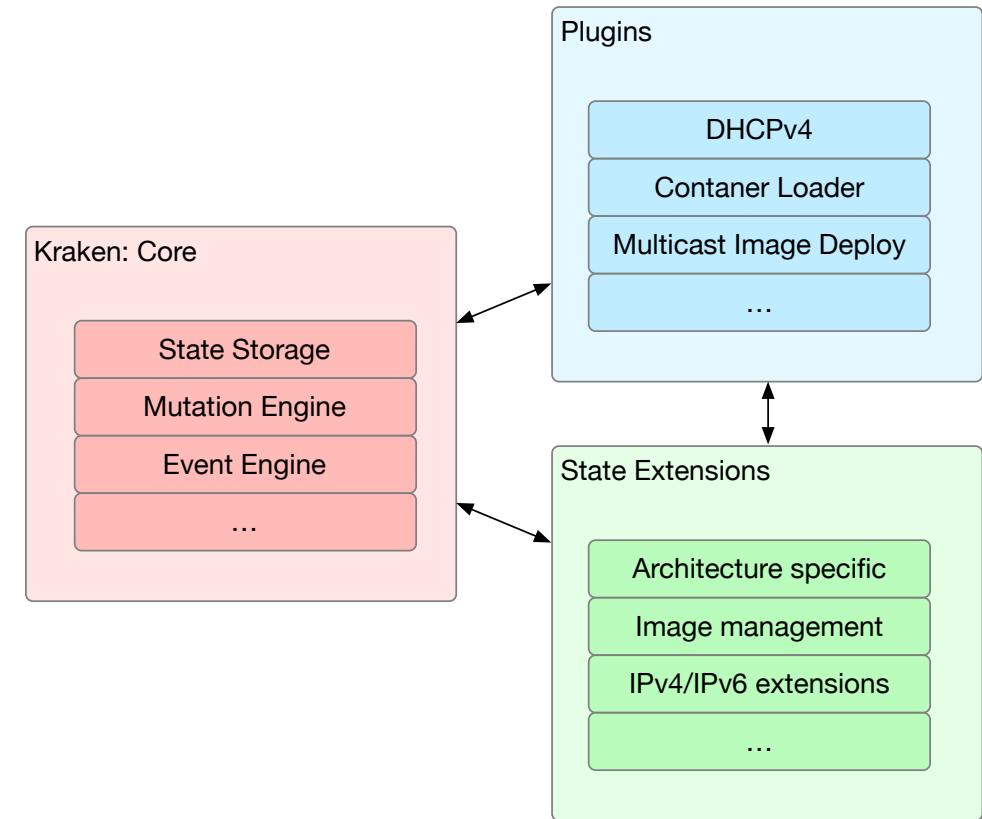
- This uses an **eventual consistency** model to converge on synchronized state.
- Discoverable state migrates **up the tree** from the endpoint nodes (computes).
- Configuration state migrates **down the tree** from the master, or “full-state node.”



Kraken is Modular

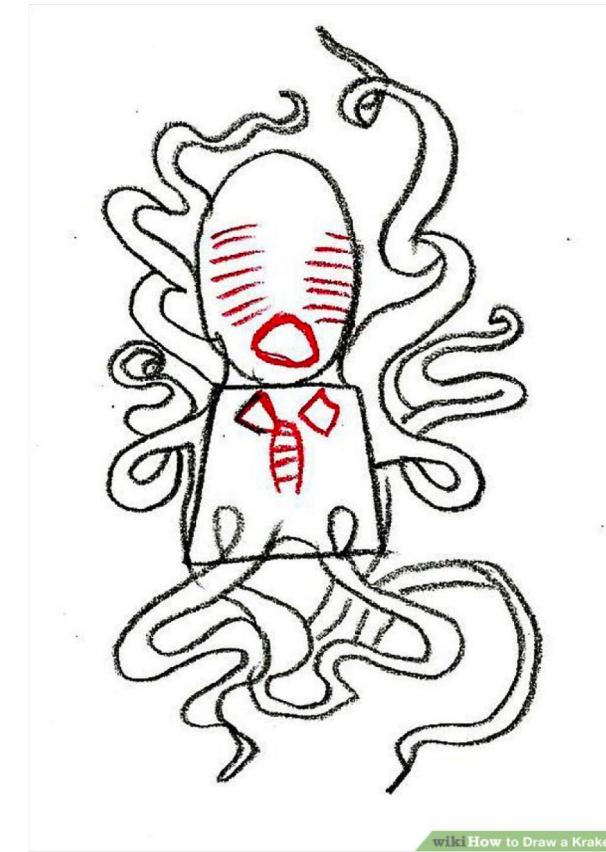
Modules make kraken actually do work.

- The **core** of Kraken:
 - State management
 - Plugin/Extension management
- **Modules** can:
 - Instantiate **microservices**
 - Declare that they can **mutate** state
 - Declare that they can **discover** state
- Extensions add new **variables** to the state



State of the Kraken

- Basic set of functional microservices & plugins
- Can boot multiple architectures
- Uses layered container images in reference implementation
- Open source (BSD-3):
<https://github.com/hpc/kraken>
- Working towards first full release



wikiHow to Draw a Kraken

Thanks!